

Regularization, Data Augmentation and Self-Supervised Learning

Efficient Deep Learning - Session 2



2023

2023-04-03

Regularization, DA and SSL

Regularization, Data Augmentation and Self-Supervised Learning

Efficient Deep Learning - Session 2



2023

└ Course organisation

Sessions

- 1 Intro Deep Learning,
- 2 Data Augmentation and Self Supervised Learning,
- 3 Quantization,
- 4 Pruning,
- 5 Factorization,
- 6 Distillation,
- 7 Embedded Software and Hardware for DL,
- 8 Presentations for challenge.

Sessions

- 1 Intro Deep Learning,
- 2 Data Augmentation and Self Supervised Learning,
- 3 Quantization,
- 4 Pruning,
- 5 Factorization,
- 6 Distillation,
- 7 Embedded Software and Hardware for DL,
- 8 Presentations for challenge.

└ Course organisation

Sessions

- 1 Intro Deep Learning,
- 2 Data Augmentation and Self Supervised Learning,
- 3 Quantization,
- 4 Pruning,
- 5 Factorization,
- 6 Distillation,
- 7 Embedded Software and Hardware for DL,
- 8 Presentations for challenge.

Sessions

- 1 Intro Deep Learning,
- 2 Data Augmentation and Self Supervised Learning,
- 3 Quantization,
- 4 Pruning,
- 5 Factorization,
- 6 Distillation,
- 7 Embedded Software and Hardware for DL,
- 8 Presentations for challenge.

Why this session ?

Regularization

Constrain the training for faster convergence and better generalization.

Data Augmentation (DA)

Help generalization by sampling training examples from a larger distribution using randomized transforms.

Self-supervised Learning (SSL)

Exploit DA and regularization tricks for learning representations, without labels

Significance

- In some (most?) cases, DA regularizes training and is needed.
- Large networks can't be trained without regularization.

2023-04-03

Regularization, DA and SSL

└ Why this session ?

Regularization prevents overfitting in neural networks, thus improve the predictions on data outside the training set

Why this session ?

Regularization

Constrain the training for faster convergence and better generalization.

Data Augmentation (DA)

Help generalization by sampling training examples from a larger distribution using randomized transforms.

Self-supervised Learning (SSL)

Exploit DA and regularization tricks for learning representations, without labels

Significance

- In some (most?) cases, DA regularizes training and is needed.
- Large networks can't be trained without regularization.

Why this session ?

Regularization

Constrain the training for faster convergence and better generalization.

Data Augmentation (DA)

Help generalization by sampling training examples from a larger distribution using randomized transforms.

Self-supervised Learning (SSL)

Exploit DA and regularization tricks for learning representations, without labels

Significance

- In some (most?) cases, DA regularizes training and is needed.
- Large networks can't be trained without regularization.

2023-04-03

Regularization, DA and SSL

└ Why this session ?

DA is a technique that increases the training set by creating new data points based on the original data during training

Why this session ?
Regularization Constrain the training for faster convergence and better generalization.
Data Augmentation (DA) Help generalization by sampling training examples from a larger distribution using randomized transforms.
Self-supervised Learning (SSL) Exploit DA and regularization tricks for learning representations, without labels
Significance <ul style="list-style-type: none">■ In some (most?) cases, DA regularizes training and is needed.■ Large networks can't be trained without regularization.

Why this session ?

Regularization

Constrain the training for faster convergence and better generalization.

Data Augmentation (DA)

Help generalization by sampling training examples from a larger distribution using randomized transforms.

Self-supervised Learning (SSL)

Exploit DA and regularization tricks for learning representations, without labels

Significance

- In some (most?) cases, DA regularizes training and is needed.
- Large networks can't be trained without regularization.

2023-04-03

Regularization, DA and SSL

└ Why this session ?

SSL in deep learning is a technique for learning data representation without the use of label for transfer learning or fine-tuning

Why this session ?
Regularization Constrain the training for faster convergence and better generalization.
Data Augmentation (DA) Help generalization by sampling training examples from a larger distribution using randomized transforms.
Self-supervised Learning (SSL) Exploit DA and regularization tricks for learning representations, without labels
Significance <ul style="list-style-type: none">■ In some (most?) cases, DA regularizes training and is needed.■ Large networks can't be trained without regularization.

Why this session ?

Regularization

Constrain the training for faster convergence and better generalization.

Data Augmentation (DA)

Help generalization by sampling training examples from a larger distribution using randomized transforms.

Self-supervised Learning (SSL)

Exploit DA and regularization tricks for learning representations, without labels

Significance

- In some (most?) cases, DA regularizes training and is needed.
- Large networks can't be trained without regularization.

2023-04-03

Regularization, DA and SSL

└ Why this session ?

Why this session ?

Regularization

Constrain the training for faster convergence and better generalization.

Data Augmentation (DA)

Help generalization by sampling training examples from a larger distribution using randomized transforms.

Self-supervised Learning (SSL)

Exploit DA and regularization tricks for learning representations, without labels

Significance

- In some (most?) cases, DA regularizes training and is needed.
- Large networks can't be trained without regularization.

Weight Decay

An old idea (Krogh and Herz 1991): ℓ_2 penalty term is added to the loss, limits the growth of model weights.

Has been shown to increase generalization and suppresses irrelevant model weights.

Ressources :

- <https://proceedings.neurips.cc/paper/1991/file/8eefcfd5990e441f0fb6f3fad709e21-Paper.pdf>
- https://ja.d2l.ai/chapter_deep-learning-basics/weight-decay.html
- Readily available in pytorch (optimizer options)

Regularization

Weight decay involves adding a term to the objective function that is proportional to the sum of the squares of the weights

Weight Decay

An old idea (Krogh and Herz 1991): ℓ_2 penalty term is added to the loss, limits the growth of model weights.

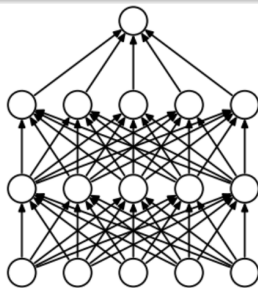
Has been shown to increase generalization and suppresses irrelevant model weights.

Ressources :

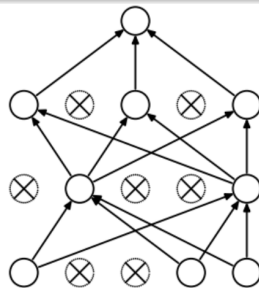
- <https://proceedings.neurips.cc/paper/1991/file/8eefcfd5990e441f0fb6f3fad709e21-Paper.pdf>
- https://ja.d2l.ai/chapter_deep-learning-basics/weight-decay.html
- Readily available in pytorch (optimizer options)

Dropout

Randomly "drops" some units during training with a certain probability.



(a) Standard Neural Net



(b) After applying dropout.

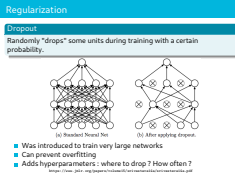
- Was introduced to train very large networks
- Can prevent overfitting
- Adds hyperparameters : where to drop ? How often ?

<https://www.jmlr.org/papers/volume15/srivastava14a/srivastava14a.pdf>

2023-04-03

Regularization, DA and SSL

└ Regularization



Dropout is a technique used only for training so that neurons will not learn redundant information of data, as well as not relying on some specific features as they might be randomly dropped out

Batch Normalization (Ioffe & Szegedy, 2015)

Normalize feature distributions to the standard distribution by learning batch statistics.

- Consider a batch X
- Calculate $m = E(X)$ and $\sigma = \text{Var}(X)$
- Compute $\hat{X} = \frac{X-m}{\sigma} * \gamma + \beta$
- m and σ are continuously updated across batches using running statistics, and γ and β are learnable parameters (by default set to 1 and 0, respectively)

Notes

- Has been shown to accelerate training, increase generalization
- Can remove the need for DropOut
- Should be included by default after convolutions

└ Regularization

- Consider a batch X
- Calculate $m = E(X)$ and $\sigma = \text{Var}(X)$
- Compute $\hat{X} = \frac{X-m}{\sigma} * \gamma + \beta$
- m and σ are continuously updated across batches using running statistics, and γ and β are learnable parameters (by default set to 1 and 0, respectively)

- Has been shown to accelerate training, increase generalization
- Can remove the need for DropOut
- Should be included by default after convolutions

Batch Normalization serves to speed up convergence, and also allows the use of higher learning rates without risk of divergence

Data Augmentation using image transformations

Translations, rotations, Scaling, Shifting in RGB, Crops,



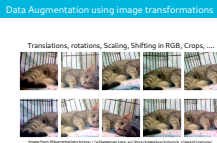
Image from Albumentations https://albumentations.ai/docs/examples/pytorch_classification/

2023-04-03

Regularization, DA and SSL

└ Data Augmentation using image transformations

DA increases the amount of data that the model sees during training, it is only applied on the training set. Note that it's stochastic meaning that the model sees different augmented versions of the images in each epoch



Mixup, Cutout and Cutmix

Mixup

For a network F trained using Cross Entropy (CE),

- Sample x_i, x_j from the training data, associated to labels y_i, y_j .
- Defined mixed up data samples as $\tilde{x} = \lambda x_i + (1 - \lambda)x_j$
- $loss = \lambda CE(F(\tilde{x}), y_i) + (1 - \lambda)CE(F(\tilde{x}), y_j)$, where $\lambda \in [0, 1]$
- Train with backprop

Notes

- Has been shown to regularize training and achieves better generalization.
- Should be included most of the time when training classification networks !
- See Lab4.md for a proposed implementation

<https://arxiv.org/pdf/1710.09412.pdf>

2023-04-03

Regularization, DA and SSL

└ Mixup, Cutout and Cutmix

Mixup, Cutout and Cutmix

Mixup

For a network F trained using Cross Entropy (CE),

- Sample x_i, x_j from the training data, associated to labels y_i, y_j .
- Defined mixed up data samples as $\tilde{x} = \lambda x_i + (1 - \lambda)x_j$
- $loss = \lambda CE(F(\tilde{x}), y_i) + (1 - \lambda)CE(F(\tilde{x}), y_j)$, where $\lambda \in [0, 1]$
- Train with backprop

Notes

- Has been shown to regularize training and achieves better generalization.
- Should be included most of the time when training classification networks !
- See Lab4.md for a proposed implementation

<https://arxiv.org/pdf/1710.09412.pdf>

Mixup forces simple linear behavior in-between training samples, it is also robust against noisy labels

Mixup, Cutout and Cutmix





Image	ResNet-50	Mixup [47]	Cutout [3]	CutMix
				
Label	Dog 1.0	Dog 0.5 Cat 0.5	Dog 1.0	Dog 0.6 Cat 0.4
ImageNet Cls (%)	76.3 (+0.0)	77.4 (+1.1)	77.1 (+0.8)	78.6 (+2.3)
ImageNet Loc (%)	46.3 (+0.0)	45.8 (-0.5)	46.7 (+0.4)	47.3 (+1.0)
Pascal VOC Det (mAP)	75.6 (+0.0)	73.9 (-1.7)	75.1 (-0.5)	76.7 (+1.1)


Table 1: Overview of the results of Mixup, Cutout, and our CutMix on ImageNet classification, ImageNet localization, and Pascal VOC 07 detection (transfer learning with SSD [23] finetuning) tasks. Note that CutMix significantly improves the performance on various tasks.

https://openaccess.thecvf.com/content_ICCV_2019/papers/Yun_CutMix_Regularization_Strategy_to_Train_Strong_Classifiers_With_Localizable_Features_ICCV_2019_paper.pdf

2023-04-03 Regularization, DA and SSL

└ Mixup, Cutout and Cutmix

Mixup boosts performance for classification (CLS) problems but degrades results for localization (LOCC) and objet detection (OO) problems. Cutout improves results for both CLS and LOC tasks but degrades performance for OO. CutMix improves results for all three tasks



Application to Self supervised Learning

Self-Supervised Learning
Learn representations of input samples without labels or annotations

How ?
Train encoders (e.g. ResNet) on pre-text tasks:
■ Self-Prediction
■ Contrastive Learning

Trained encoders are expected to learn general features that generalize to supervised tasks.

Self-Supervised Learning

Learn representations of input samples without labels or annotations

How ?


Train encoders (e.g. ResNet) on pre-text tasks:

- Self-Prediction
- Contrastive Learning

Trained encoders are expected to learn general features that generalize to supervised tasks.

SSL consists in learning a model that outputs useful representations of data without the use of label. This model can be used for transfer learning to other tasks. The model is trained in a supervised way with "labels" created from the data itself

Self-Prediction : Colorful Image Colorization



<https://arxiv.org/pdf/1603.08511.pdf>

Colorful Image Colorization


Learn feature representations by restoring colored version of images :

- Given lightness L, predict a and b color channels (CIE Lab colorspace)
- A loss penalty is computed between the predicted and the original image

2023-04-03

Regularization, DA and SSL

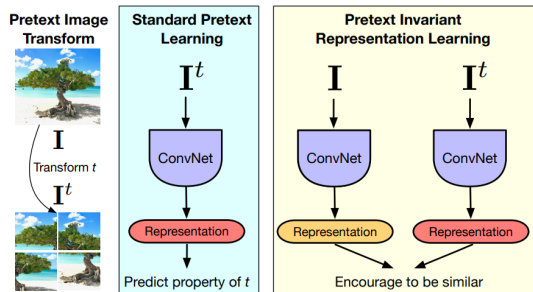
Application to Self supervised Learning



Colorizing images is an example of self-prediction tasks that learns representation of data by predicting a part of the data (colors in this case)

Application to Self supervised Learning

Contrastive Learning : Pretext-Invariant Representations Learning (PIRL)



<https://arxiv.org/pdf/1912.01991.pdf>

PIRL


Train a discriminative feature extractor :

- Images and their augmented version should have similar representations
- Different images should have dissimilar representations

2023-04-03

Regularization, DA and SSL

└ Application to Self supervised Learning



Contrastive learning is the state-of-the-art method for learning good representation of data, it consists in learning discriminative features by measuring similarities and dissimilarities between samples, it relies heavily on data augmentation