Quantizing neural networks

Efficient Deep Learning - Session 3

IMT Atlantique
Bretagne-Pays de la Loire
École Mines-Télécom

Course organisation

Sessions

1. Intro Deep Learning,
2. Data Augmentation and Self Supervised Learning,
3. Quantization,
4. Pruning,
5. Factorization,
6. Distillation,
7. Embedded Software and Hardware for DL,
8. Presentations for challenge.

Course organisation

Sessions

1. Intro Deep Learning,
2. Data Augmentation and Self Supervised Learning,
3. Quantization,
4. Pruning,
5. Factorization,
6. Distillation,
7. Embedded Software and Hardware for DL,
8. Presentations for challenge.

Today's Summary
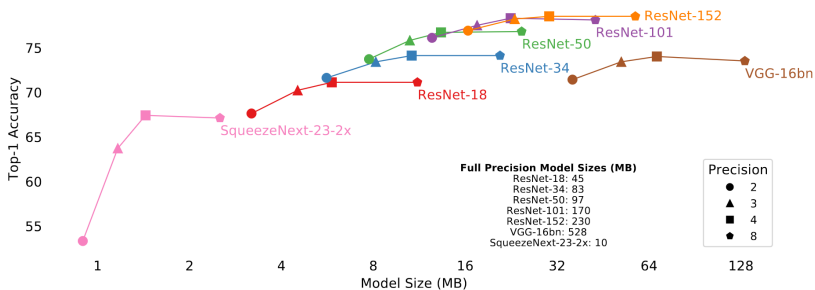
Plan

Motivation

1 Reduce model size
  - Fewer bits → Reduced memory footprint

2 Decrease memory access
  - GPU & CPU : reduce Cache usage

3 Computational complexity

1. Reduce model size
   - Fewer bits $\rightarrow$ Reduced memory footprint
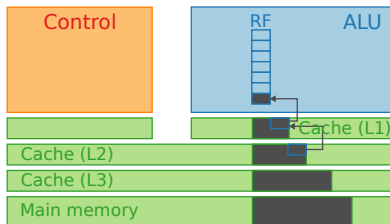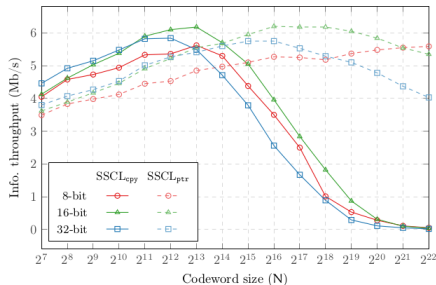2. Decrease memory access
   - GPU & CPU : reduce Cache usage
3. Computational complexity

Motivation

1. Reduce model size
   - Fewer bits $\rightarrow$ Reduced memory footprint
2. Decrease memory access
   - GPU & CPU : reduce Cache usage
3. Computational complexity

Motivation

1. **Reduce model size**
   - Fewer bits $\rightarrow$ Reduced memory footprint
2. **Decrease memory access**
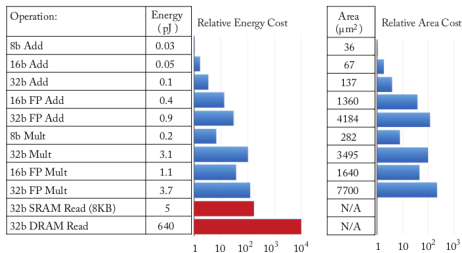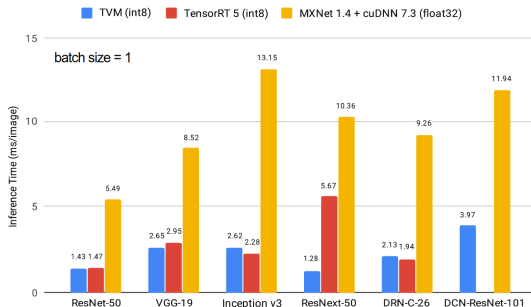   - GPU & CPU : reduce Cache usage
3. **Computational complexity**

Figure 7.1: The area and energy cost for additions and multiplications at different precision, and memory accesses in a 45 nm process. The area and energy scale different for multiplication and addition. The energy consumption of data movement (red) is significantly higher than arithmetic operations (blue). (Figure adapted from [121].)

From : Sze, Vivienne, et al. "Efficient processing of deep neural networks." Synthesis Lectures on Computer Architecture

Motivation

1 Reduce model size
- Fewer bits → Reduced memory footprint

2 Decrease memory access
- GPU & CPU : reduce Cache usage

3 Computational complexity

From : https://github.com/vinx13/tvm-cuda-int8-benchmark/
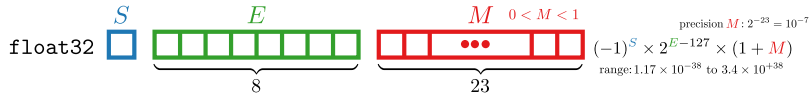
Plan

Floating Point

Floating Point

Floating Point

float32

S, E (8), M (23), $0 < M < 1$

precision $M: 2^{-23} = 10^{-7}$

$(-1)^S \times 2^{E-127} \times (1 + M)$

range: 1.17×10^{-38} to $3.4 \times 10^{+38}$

float16

S, E (5), M (10), $0 < M < 1$

precision $M: 2^{-10} \approx 0.001$

$(-1)^S \times 2^{E-15} \times (1 + M)$

range: 3×10^{-5} to $1.3 \times 10^{+5}$

bfloat16

S, E (8), M (7), $0 < M < 1$

precision $M: 2^{-7} \approx 0.001$

$(-1)^S \times 2^{E-127} \times (1 + M)$

range: 1.17×10^{-38} to $3.4 \times 10^{+38}$

float32 — S | E (8) | M ($0 < M < 1$) (23)

$(-1)^{S} \times 2^{E-127} \times (1 + M)$

precision $M : 2^{-23} = 10^{-7}$

range: 1.17×10^{-38} to $3.4 \times 10^{+38}$

float16 — S | E (5) | M ($0 < M < 1$) (10)

$(-1)^{S} \times 2^{E-15} \times (1 + M)$

precision $M : 2^{-10} \approx 0.001$

range: 3×10^{-5} to $1.3 \times 10^{+5}$

bfloat16 — S | E (8) | M ($0 < M < 1$) (7)

$(-1)^{S} \times 2^{E-127} \times (1 + M)$

precision $M : 2^{-7} \approx 0.001$

range: 1.17×10^{-38} to $3.4 \times 10^{+38}$
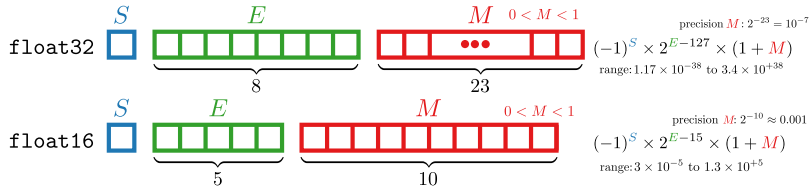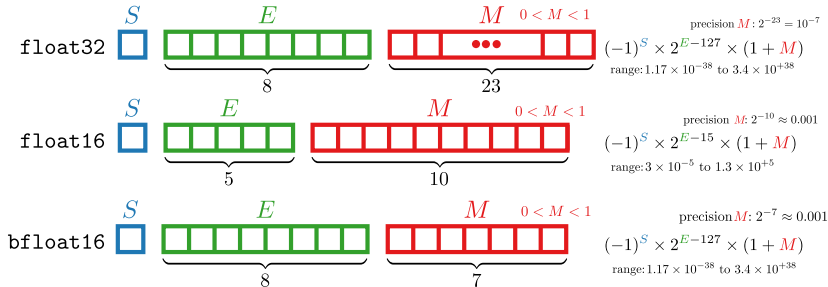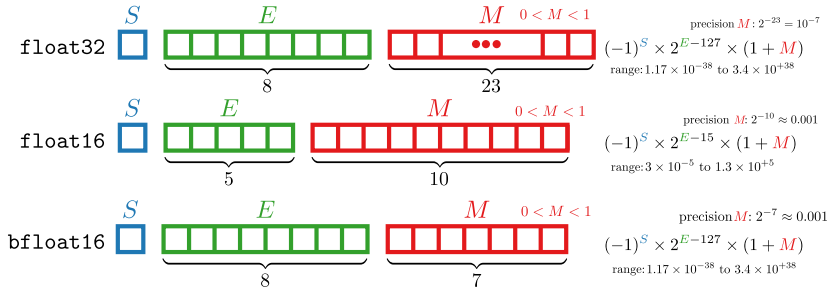
- To add two FP numbers:
  - Shift M according to E (int shift n_E bits)
  - Add M (int add n_M bits)
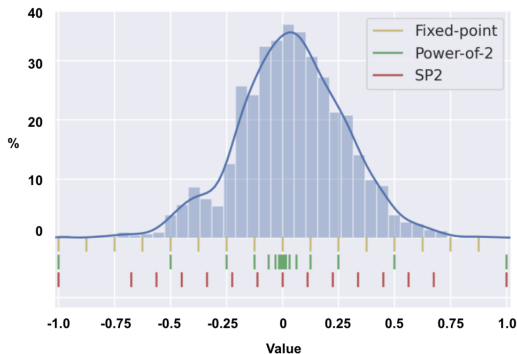 - Normalize ($0 < M < 1$)
- To multiply two FP numbers:
  - Multiply M (int mult n_M bits)
  - Add E (int mult n_E bits)
  - Normalize ($0 < M < 1$)

Integers, fixed point

int32 with 32 bits, $\times 2^{-i}$, $0 \rightarrow 2^{32} - 1 \approx 4.3 \times 10^6$

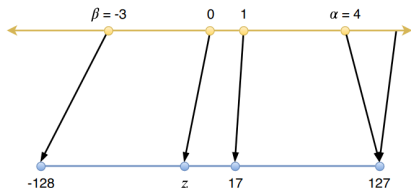int8 with 8 bits, $\times 2^{-i}$, $0 \rightarrow 255$
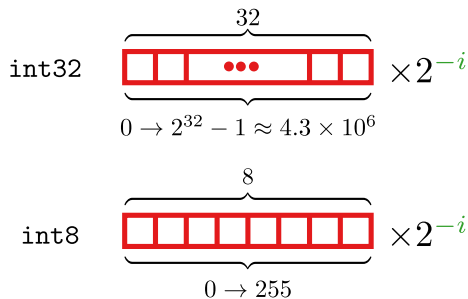
- Fixed point ($-i$)
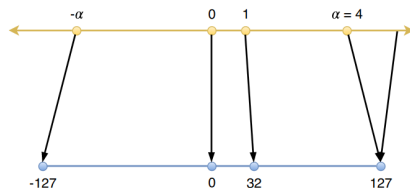- Short range
- Simple computation

Uniform and Non-Uniform Quantization

- Uniform quantization enables the use of integer on fixed-point hardware
- Non-uniform quantization requires a codebook lookup $\rightarrow$ not straightforward for standard hardware (CPU, GPU)

Affine and Scale Quantization
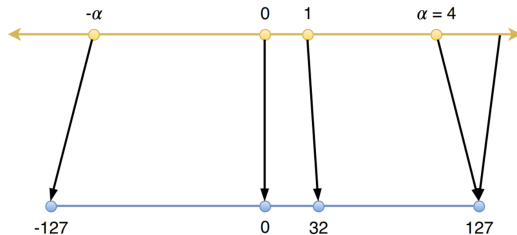

(a) Affine quantization

(b) Scale quantization

- 2 kinds of uniform quantization
- Assymetric vs Symmetric

Wu, Hao, et al. "Integer quantization for deep learning inference: Principles and empirical evaluation." arXiv preprint arXiv:2004.09602 (2020).

Scale Quantization

$$\text{clip}(x, l, u) \begin{cases} l, & x < l \\ x, & l \leq x \leq u \\ u, & x > u \end{cases}$$

$$s = \frac{2^{b-1} - 1}{\alpha}$$

$$x_q = \text{quantize}(x, b, s) = \text{clip}(\text{round}(s \cdot x), -2^{b-1} + 1, 2^{b-1} - 1)$$

$$\hat{x} = \text{dequantize}(x_q, s) = \frac{1}{s} x_q$$

Scale Quantization

$$y_{ij} = \sum_{k=1}^{p} x_{ik} \cdot w_{kj} \approx$$

$$\sum_{k=1}^{p} \mathrm{dequantize}(x_{q,ik}, s_{q,ik}) \cdot \mathrm{dequantize}(w_{q,kj}, s_{w,kj}) =$$

$$\sum_{k=1}^{p} \frac{1}{s_{x,ik}} x_{q,ik} \cdot \frac{1}{s_{w,kj}} w_{q,kj}$$

And, in order to use integer multiplication, the scaling factor s must be independent of k :

$$\frac{1}{s_{x,i} \cdot s_{w,j}} \sum_{k=1}^{p} x_{q,ik} \cdot w_{q,kj}$$

Plan

Quantization Post Training : Weights

Quantize a set of parameters

- Select the number of quantization bits,
- Select the set of parameters to quantize,
- Determine range of values,
- Determine the scaling factor (and zero if affine).

Different weight sets can be considered

- Whole network,
- per layer,
- per neuron.

Finer sets segmentation $\rightarrow$ better accuracy.

Quantization Post Training : Activation

Quantize a set of parameters

- Select the number of quantization bits,
- Select the se t of activations to quantize,
- Determine range of values **according to the training set or a subset**,
- Determine the scaling factor (and zero if affine).

Different weight sets can be considered

- Whole network,
- per layer,
- per neuron.

Finer sets segmentation $\rightarrow$ better accuracy.

Quantization Aware Training

- Train the network while quantizing the weights and / or the activations,
- Quantization Aware Techniques yield way better accuracy,
- Especially for extremely low-bit precision (2-3-4 bit precision).

Learned Step Size Quantization
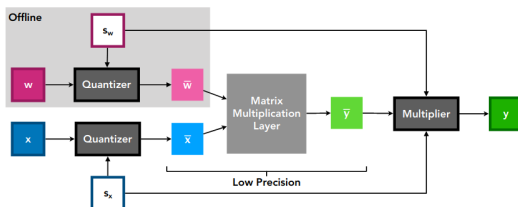
- s quantizer step size
- Q_P and Q_N, the number of positive and negative quantization levels

$$\bar{v} = \lfloor clip(v/s, -Q_N, Q_P) \rceil, \tag{1}$$

$$\hat{v} = \bar{v} \times s. \tag{2}$$

- s is learned with :

$$\frac{\partial \hat{v}}{\partial s} = \begin{cases} -v/s + \lfloor v/s \rceil & \text{if } -Q_N < v/s < Q_P \\ -Q_N & \text{if } v/s \leq -Q_N \\ Q_P & \text{if } v/s \geq Q_P \end{cases} \tag{3}$$

Learned Step Size Quantization - `https://arxiv.org/pdf/1902.08153.pdf`

Quantization while Learning - Binary Connect

Algorithm 1 SGD training with BinaryConnect. C is the cost function for minibatch and the functions binarize(w) and clip(w) specify how to binarize and clip weights. L is the number of layers.

Require: a minibatch of (inputs, targets), previous parameters w_{t-1} (weights) and b_{t-1} (biases), and learning rate η.

Ensure: updated parameters w_t and b_t.

  1. Forward propagation:

  $w_b \leftarrow$ binarize(w_{t-1})

  For $k = 1$ to L, compute a_k knowing a_{k-1}, w_b and b_{t-1}

  2. Backward propagation:

  Initialize output layer's activations gradient $\frac{\partial C}{\partial a_L}$

  For $k = L$ to 2, compute $\frac{\partial C}{\partial a_{k-1}}$ knowing $\frac{\partial C}{\partial a_k}$ and w_b

  3. Parameter update:

  Compute $\frac{\partial C}{\partial w_b}$ and $\frac{\partial C}{\partial b_{t-1}}$ knowing $\frac{\partial C}{\partial a_k}$ and a_{k-1}

  $w_t \leftarrow$ clip($w_{t-1} - \eta \frac{\partial C}{\partial w_b}$)

  $b_t \leftarrow b_{t-1} - \eta \frac{\partial C}{\partial b_{t-1}}$

   Courbariaux, Matthieu, Yoshua Bengio, and Jean-Pierre David. "Binaryconnect: Training deep neural networks with binary weights during propagations." Advances in neural information processing systems. 2015.
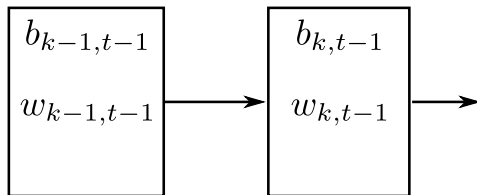`https://arxiv.org/pdf/1511.00363.pdf`

1. Forward propagation:

$w_b \leftarrow \text{binarize}(w_{t-1})$

For $k = 1$ to L, compute a_k knowing a_{k-1}, w_b and b_{t-1}
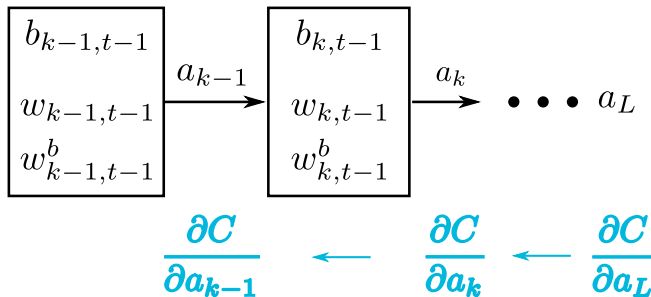
1. Forward propagation:

$w_b \leftarrow \text{binarize}(w_{t-1})$

$$w^b = \begin{cases} +1, & \text{if } w \geq 0 \\ -1, & \text{otherwise} \end{cases}$$

1. Forward propagation:

$w_b \leftarrow \text{binarize}(w_{t-1})$

For $k = 1$ to L, compute a_k knowing a_{k-1}, w_b and b_{t-1}

2. Backward propagation:

Initialize output layer's activations gradient $\frac{\partial C}{\partial a_L}$

For $k = L$ to 2, compute $\frac{\partial C}{\partial a_{k-1}}$ knowing $\frac{\partial C}{\partial a_k}$ and w_b

3. Parameter update:

Compute $\frac{\partial C}{\partial w_b}$ and $\frac{\partial C}{db_{t-1}}$ knowing $\frac{\partial C}{\partial a_k}$ and a_{k-1}

$w_t \leftarrow \text{clip}(w_{t-1} - \eta \frac{\partial C}{\partial w_b})$

$b_t \leftarrow b_{t-1} - \eta \frac{\partial C}{\partial b_{t-1}}$

- Deterministic

$$w_b = \begin{cases} +1, & \text{if } w \geq 0 \\ -1, & \text{otherwise} \end{cases}$$

- Stochastic

$$w_b = \begin{cases} +1, & \text{with probability } p = \sigma(w) \\ -1, & \text{with probability } 1 - p \end{cases}$$

avec

$$\sigma(x) = \text{clip}(\frac{x+1}{2}, 0, 1) = \max(0, \min(1, \frac{x+1}{2}))$$

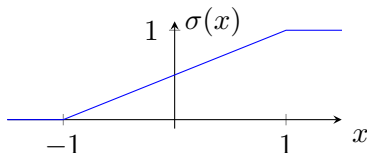# Binarization : Stochastic vs Deterministic

- Deterministic

$$w_b = \begin{cases} +1, & \text{if } w \geq 0 \\ -1, & \text{otherwise} \end{cases}$$

- Stochastic

$$w_b = \begin{cases} +1, & \text{with probability } p = \sigma(w) \\ -1, & \text{with probability } 1 - p \end{cases}$$

avec

$$\sigma(x) = \text{clip}(\frac{x+1}{2}, 0, 1) = \max(0, \min(1, \frac{x+1}{2}))$$

Plan

Quantization in Pytorch

1. Dynamic Quantization
2. Static Quantization
3. Quantization Aware Training

https://pytorch.org/blog/introduction-to-quantization-on-pytorch/

And for our need : `https://pytorch.org/tutorials/prototype/fx_graph_mode_ptq_static.html`